

title: “ASI Coursework” author: “Margaret Duff” date: “30 November 2018”  
 output: word\_document: default html\_document: df\_print: paged geometry:  
 margin=0.2cm —

## Fatigue of materials - part 4

Consider now the following random effects model that allows the fatigue limit to be different for each coupon. This is achieved by modelling the fatigue limit as an unobserved random variable  $\Gamma$ . For coupon  $i$ , the conditional distribution of the number of cycles to failure  $N_i$  given that  $\Gamma_i = \gamma_i < s_i$  that is, given that the realization of the fatigue limit for that coupon is below the applied stress level is given by

$$N_i | \Gamma_i = \gamma_i < s_i \sim \text{Weibull}(\text{shape} = 1/\sigma, \text{scale} = \alpha (s_i - \gamma_i)^\delta) \quad (1)$$

We will assume that  $\Gamma_1, \dots, \Gamma_{26}$  are i.i.d:

$$\text{Weibull}(\text{shape} = 1/\sigma_\gamma, \text{scale} = \exp(\mu_\gamma)) \quad (2)$$

where  $\mu_\gamma \in R, \sigma_\gamma > 0$  are unknown parameters. Let  $b = (\gamma_1, \dots, \gamma_{26})^\top$  and the vector of unknown parameters is now given by  $\theta^\top = (\log(\alpha), \delta, \log(\sigma), \mu_\gamma, \log(\sigma_\gamma))$ . We will assume the following priors:  $\log(\alpha), \delta, \log(\sigma), \mu_\gamma$  are independent and with improper uniform priors while  $\sigma_\gamma$  is exponential with rate 5 and independent of  $\log(\alpha), \delta, \log(\sigma)$  and  $\mu_\gamma$ .

We wish to use a Metropolis-Hastings algorithm to sample from the posterior distribution of  $\theta$  and the random effects  $b$ . For an initial step we split the data set into those in which the coupon broke and those which didn't.

```
split_fatigue <- split(fatigue, fatigue$ro)
broke=as.data.frame(split_fatigue[[1]])
runoff=as.data.frame(split_fatigue[[2]])
```

The below computes a log posterior function for the model given parameters:  $\log(\alpha), \delta, \log(\sigma), \mu_\gamma, \log(\sigma_\gamma)$  and  $\gamma$

```
# function to compute log posterior
log.post <- function(log_alpha, delta, log_sigma, mu_gamma, log_sigma_gamma, gamma_broke, gamma_runoff) {
  log_pi0_log_alpha=log(1);log_pi0_delta= log(1);log_pi0_log_sigma=log(1);log_pi0_mu_gamma=log(1)
  log_pi0_log_sigma_gamma=log(5*exp(-5*exp(log_sigma_gamma)))
  log_pi0_gamma_broke=log(dweibull(gamma_broke, 1/exp(log_sigma_gamma), exp(mu_gamma)))#log
  log_pi0_gamma_runoff=log(dweibull(gamma_runoff, 1/exp(log_sigma_gamma), exp(mu_gamma)))#log
  log_pi0<- log_pi0_log_alpha+log_pi0_delta+log_pi0_log_sigma+log_pi0_mu_gamma+log_pi0_log_sigma_gamma+log_pi0_gamma_broke+log_pi0_gamma_runoff
  log.lik=sum(log(pweibull(broke$N, 1/exp(log_sigma), exp(log_alpha)*(broke$s-gamma_broke)^(delta), log_sigma_gamma)))
  return(log.lik+log_pi0) # now the log posterior = log likelihood +log prior
}
```

We now move onto the MH algorithm. for the proposal distributions we consider a normal random variable centered on the previous value with variance that we can tune for the variables  $\log(\alpha)$ ,  $\delta$ ,  $\log(\sigma)$ ,  $\mu_\gamma$  and  $\log(\sigma_\gamma)$  and a uniform random variable centered on the previous value with range that we can vary for  $\gamma$ . We choose a burn in period of 10000 although this can be altered and tune the parameters to try and get acceptance rates of around 25% although as noted earlier, the system is quite sensitive to values of  $\gamma$  making the tuning very difficult to do.

```
nsteps=100000; burn.in=10000
MH<-function(log_alpha0, delta0, log_sigma0,mu_gamma0,log_sigma_gamma0, gamma_broke0, gamma_runoff0){
  accept <- rep(0,3)# will keep track of acceptance rates for the inner middle and outer me
  log_alpha <- rep(0,nsteps); delta=rep(0,nsteps);log_sigma=rep(0,nsteps); mu_gamma=rep(0,nsteps)
  gamma_broke=matrix(0, nsteps, length(broke$N)); gamma_runoff=matrix(0, nsteps, length(runoff$N))
  log_alpha[1] <- log_alpha0; delta[1]=delta0; log_sigma[1]=log_sigma0; mu_gamma[1]=mu_gamma0
  gamma_broke[1,]=gamma_broke0; gamma_runoff[1,]=gamma_runoff0#Set initial values
  lp0 <- log.post(log_alpha0, delta0, log_sigma0,mu_gamma0,log_sigma_gamma0, gamma_broke0, gamma_runoff0)
  for (i in 2:nsteps){
    #Set current values
    current_log_alpha=log_alpha[i-1]; current_delta=delta[i-1]; current_log_sigma=log_sigma[i-1]
    current_gamma_broke=gamma_broke[i-1,];current_gamma_runoff=gamma_runoff[i-1,] #extract current values
    proposed_mu_gamma=current_mu_gamma+rnorm(1,0, range_mu_gamma) # new proposed values
    proposed_log_sigma_gamma=current_log_sigma_gamma+rnorm(1,0, range_log_sigma_gamma) # new proposed values
    lp1 <- log.post(current_log_alpha, current_delta, current_log_sigma,proposed_mu_gamma,proposed_log_sigma_gamma,
      current_gamma_broke,current_gamma_runoff)
    acc <- exp(min(0,lp1-lp0))
    if (runif(1) >= acc | !is.finite(acc)){# reject
      log_alpha[i] <- current_log_alpha; delta[i]=current_delta; log_sigma[i]=current_log_sigma
      lp1=lp0 # keep log posterior values up to date
    }else{#accept
      accept[1]=accept[1]+1 # keep track to calculate acceptance rates
      mu_gamma[i]=proposed_mu_gamma ; log_sigma_gamma[i]=proposed_log_sigma_gamma # update values
      lp0=lp1# keep log posterior values up to date
      proposed_gamma_broke= current_gamma_broke + runif(length(broke$N), -range_gamma, range_gamma)
      proposed_gamma_runoff= current_gamma_runoff+ runif(length(runoff$N), -range_gamma, range_gamma)
      check1=isTRUE(all.equal(abs(broke$s-proposed_gamma_broke), broke$s-proposed_gamma_broke))
      check2=isTRUE(all.equal(abs(runoff$s-proposed_gamma_runoff), runoff$s-proposed_gamma_runoff))
      if(!(check1 & check2) ){ # reject
        log_alpha[i] <- current_log_alpha;delta[i]=current_delta;log_sigma[i]=current_log_sigma
        lp1=lp0# keep log posterior values up to date
      } else{# continue with MH
        lp1 <- log.post(current_log_alpha, current_delta, current_log_sigma,proposed_mu_gamma,proposed_log_sigma_gamma,
          proposed_gamma_broke,proposed_gamma_runoff)
        acc <- exp(min(0,lp1-lp0))
        if (runif(1) >= acc | !is.finite(acc)){# reject
          log_alpha[i] <- current_log_alpha;delta[i]=current_delta; log_sigma[i]=current_log_sigma
          lp1=lp0# keep log posterior values up to date
        }else{ #accept
```

```

        accept[2]=accept[2]+1 # keep track to calculate acceptance rates
        gamma_broke[i,]=proposed_gamma_broke; gamma_runoff[i,]=proposed_gamma_runoff # update
        lp0=lp1# keep log posterior values up to date
        proposed_log_alpha=current_log_alpha+rnorm(1, 0, range_log_alpha) #new proposed values
        proposed_delta=current_delta+rnorm(1, 0, range_delta) #new proposed values
        proposed_log_sigma=current_log_sigma+rnorm(1, 0, range_log_sigma) #new proposed values
        lp1 <- log.post(proposed_log_alpha, proposed_delta, proposed_log_sigma,proposed_mu_gamma)
        acc <- exp(min(0,lp1-lp0))
        if (runif(1) >= acc| !is.finite(acc)){# reject
            log_alpha[i] <- current_log_alpha; delta[i]=current_delta;log_sigma[i]=current_log_sigma
            lp1=lp0# keep log posterior values up to date
        }else{#accept
            accept[3]=accept[3]+1 # keep track to calculate acceptance rates
            log_alpha[i] <- proposed_log_alpha
            delta[i]=proposed_delta ; log_sigma[i]=proposed_log_sigma # keep track of variables
            lp0=lp1# keep log posterior values up to date
        } } } }
list(log_alpha=log_alpha, delta=delta, log_sigma=log_sigma,mu_gamma=mu_gamma,log_sigma_gamma=log_sigma_gamma)
}
mh=MH(0,1,1,4,-3,rep(50, length(broke$N)),rep(50, length(runoff$N)),0.1,0.3,0.3,0.05,0.05,0.05)
mh$ar_outer; mh$ar_middle; mh$ar_inner

```

```

## [1] 0.32513
## [1] 0.8808169
## [1] 0.5231161


```

We plot the trace graphs of our retained results minus the burn-in period:

```

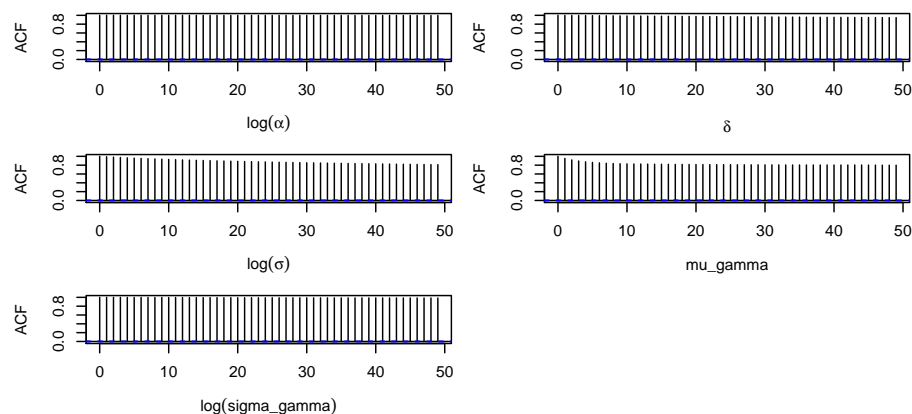
show.plot<- (burn.in):nsteps
par(mfrow=c(3,2), mar=c(4,4,1,1))
plot(mh$log_alpha[show.plot],type="l",ylab=expression(log_alpha))
plot(mh$delta[show.plot],type="l",ylab=expression(delta))
plot(mh$log_sigma[show.plot],type="l",ylab=expression(log(sigma)))
plot(mh$mu_gamma[show.plot],type="l",ylab=expression(mu_gamma))
plot(mh$log_sigma_gamma[show.plot],type="l",ylab=expression(log(sigma_gamma)))

```



and the auto-correlation functions:

```
par(mfrow=c(3,2),mar=c(4,4,1,1))
acf(mh$log_alpha[-burn.in],xlab=expression(log(alpha)))
acf(mh$delta[-burn.in],xlab=expression(delta))
acf(mh$log_sigma[-burn.in],xlab=expression(log(sigma)))
acf(mh$mu_gamma[-burn.in],xlab=expression(mu_gamma))
acf(mh$log_sigma_gamma[-burn.in],xlab=expression(log(sigma_gamma)))
```

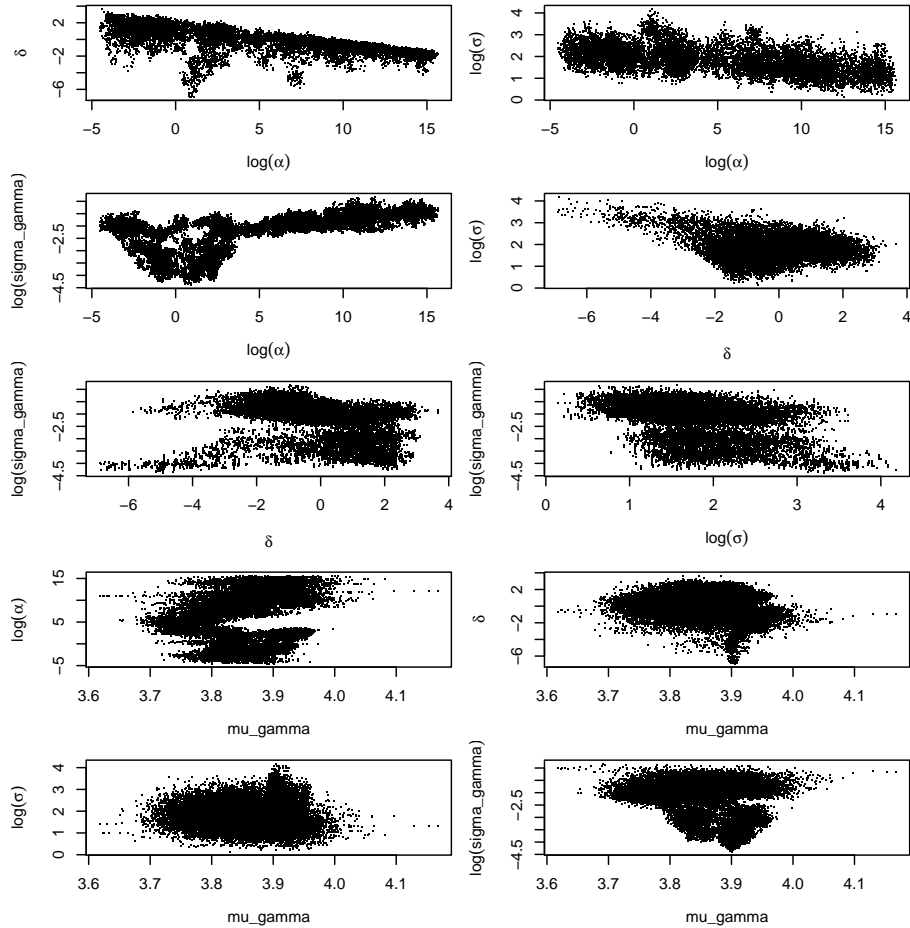


We see that we get quite small values for effective sample size and that the sample size for  $\log(\sigma_\gamma)$  is extremely small.

```
n.eff <- c(0,0,0,0,0)
autocor <- acf(mh$log_alpha[-(burn.in)],plot=FALSE);t.eff <- 2*sum(autocor[[1]]) - 1;n.eff[1] <- t.eff
autocor <- acf(mh$delta[-(burn.in)],plot=FALSE);t.eff <- 2*sum(autocor[[1]]) - 1;n.eff[2] <- t.eff
autocor <- acf(mh$log_sigma[-(burn.in)],plot=FALSE);t.eff <- 2*sum(autocor[[1]]) - 1;n.eff[3] <- t.eff
autocor <- acf(mh$mu_gamma[-(burn.in)],plot=FALSE);t.eff <- 2*sum(autocor[[1]]) - 1;n.eff[4] <- t.eff
autocor <- acf(mh$log_sigma_gamma[-(burn.in)],plot=FALSE);t.eff <- 2*sum(autocor[[1]]) - 1;n.eff[5] <- t.eff
n.eff
```

```
## [1] 1011.664 1040.275 1146.975 1018.847    0.000
```

We check for correlation between the parameters by plotting graphs of a random sample of the iterations retained after discarding the burn-in:



We can't see a strong correlation between posterior values, and certainly no nice ellipses we saw in the previous part which allowed us to use a multivariate normal centered on the previous values for our proposal. There does not seem to be an obvious way of improving the metropolis Hastings sampler. Instead we consider numerical integration.